

```
In [ ]: ## Load libraries
import pandas as pd
import numpy as np
import sys
import matplotlib.pyplot as plt
import matplotlib.cm as cm
from keras.datasets import mnist
plt.style.use('dark_background')
%matplotlib inline

WARNING:tensorflow:From c:\Users\SA RAVI\anaconda3\envs\aimlsem1\lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

In [ ]: np.set_printoptions(precision=2)

In [ ]: import tensorflow as tf

In [ ]: tf.__version__

Out[ ]: '2.15.0'
```

Load MNIST Data

```
In [ ]: ## Load MNIST data
(X_train, y_train), (X_test, y_test) = mnist.load_data()
X_train = X_train.transpose(1, 2, 0)
X_test = X_test.transpose(1, 2, 0)
X_train = X_train.reshape(X_train.shape[0]*X_train.shape[1], X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0]*X_test.shape[1], X_test.shape[2])

num_labels = len(np.unique(y_train))
num_features = X_train.shape[0]
num_samples = X_train.shape[1]

# One-hot encode class labels
Y_train = tf.keras.utils.to_categorical(y_train).T
Y_test = tf.keras.utils.to_categorical(y_test).T

# Normalize the samples (images)
xmax = np.amax(X_train)
xmin = np.amin(X_train)
X_train = (X_train - xmin) / (xmax - xmin) # all train features turn into a number between 0 and 1
X_test = (X_test - xmin) / (xmax - xmin)

print('MNIST set')
print('-----')
print('Number of training samples = %d'%(num_samples))
print('Number of features = %d'%(num_features))
print('Number of output labels = %d'%(num_labels))

MNIST set
-----
Number of training samples = 60000
Number of features = 784
Number of output labels = 10
```

A generic layer class with forward and backward methods

```
In [ ]: class Layer:
def __init__(self):
    self.input = None
    self.output = None

def forward(self, input):
    pass

def backward(self, output_gradient, learning_rate):
    pass
```

The softmax classifier steps for a batch of comprising  $b$  samples represented as the  $725 \times b$ -matrix (724 pixel values plus the bias feature absorbed as its last row)

$$\mathbf{X} = [\mathbf{x}^{(0)}, \mathbf{x}^{(1)}, \dots, \mathbf{x}^{(b-1)}]$$

with one-hot encoded true labels represented as the  $10 \times b$ -matrix (10 possible categories)

$$\mathbf{Y} = [\mathbf{y}^{(0)} \quad \dots \quad \mathbf{y}^{(b-1)}]$$

using a randomly initialized  $10 \times 725$ -weights matrix  $\mathbf{W}$ :

1. Calculate  $10 \times b$ -raw scores matrix :

$$\begin{aligned} [\mathbf{z}^{(0)} \quad \dots \quad \mathbf{z}^{(b-1)} \dots] &= \mathbf{W} [\mathbf{z}^{(0)} \quad \dots \quad \mathbf{z}^{(b-1)} \dots] \\ &= [\mathbf{W}\mathbf{z}^{(0)} \quad \dots \quad \mathbf{W}\mathbf{z}^{(b-1)}] \\ &\Rightarrow \mathbf{Z} = \mathbf{W}\mathbf{X}. \end{aligned}$$

2. Calculate  $10 \times b$ -softmax predicted probabilities matrix:

$$\begin{aligned} [\mathbf{a}^{(0)} \quad \dots \quad \mathbf{a}^{(b-1)}] &= [\text{softmax}(\mathbf{z}^{(0)}) \quad \dots \quad \text{softmax}(\mathbf{z}^{(b-1)})] \\ &\Rightarrow \mathbf{A} = \text{softmax}(\mathbf{Z}). \end{aligned}$$


3. Predicted probability matrix get a new name:  $\hat{\mathbf{Y}} = \mathbf{A}$ .  
4. The crossentropy (CCE) loss for the  $i$ th sample is

$$L_i = \sum_{k=0}^9 -y^{(i)} \log(\hat{y}_k^{(i)}) = -\mathbf{y}^{(i)\top} \log(\hat{\mathbf{y}}^{(i)})$$

which leads to the average crossentropy (CCE) batch loss for the batch as:

$$\begin{aligned} L &= \frac{1}{b} [L_0 + \dots + L_{b-1}] \\ &= \frac{1}{b} \left[ -\mathbf{y}^{(0)\top} \log(\hat{\mathbf{y}}^{(0)}) + \dots + -\mathbf{y}^{(b-1)\top} \log(\hat{\mathbf{y}}^{(b-1)}) \right]. \end{aligned}$$

5. The computational graph for the samples in the batch are presented below:



6. Calculate the gradient of the average batch loss w.r.t. weights as:

$$\Rightarrow \nabla_{\mathbf{W}}(L) = \frac{1}{b} \left( \underbrace{\left[ \nabla_{\mathbf{W}}(\mathbf{z}^{(0)}) \times \nabla_{\mathbf{z}^{(0)}}(\hat{\mathbf{y}}^{(0)}) \times \nabla_{\hat{\mathbf{y}}^{(0)}}(L_0) \right]}_{\text{sample 0}} + \dots + \underbrace{\left[ \nabla_{\mathbf{W}}(\mathbf{z}^{(b-1)}) \times \nabla_{\mathbf{z}^{(b-1)}}(\hat{\mathbf{y}}^{(b-1)}) \times \nabla_{\hat{\mathbf{y}}^{(b-1)}}(L_{b-1}) \right]}_{\text{sample } b-1} \right)$$

$$= \frac{1}{b} \left( \underbrace{\left[ \nabla_{\mathbf{W}}(\mathbf{z}^{(0)}) \times \nabla_{\mathbf{z}^{(0)}}(\mathbf{a}^{(0)}) \times \nabla_{\hat{\mathbf{y}}^{(0)}}(L_0) \right]}_{\text{sample 0}} + \dots + \underbrace{\left[ \nabla_{\mathbf{W}}(\mathbf{z}^{(b-1)}) \times \nabla_{\mathbf{z}^{(b-1)}}(\hat{\mathbf{y}}^{(b-1)}) \times \nabla_{\hat{\mathbf{y}}^{(b-1)}}(L_{b-1}) \right]}_{\text{sample } b-1} \right).$$

10. The full gradient can be written as  $\nabla_{\mathbf{W}}(L) =$

$$\nabla_{\mathbf{W}}(L) = \left\{ \sum_{i=0}^{b-1} \begin{bmatrix} \mathbf{z}^{(i)} & \mathbf{0} & \mathbf{0} & \dots & \mathbf{0} \\ \mathbf{0} & \mathbf{x}^{(i)} & \mathbf{0} & \dots & \mathbf{0} \\ \mathbf{0} & \mathbf{0} & \mathbf{x}^{(i)} & \dots & \mathbf{0} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ \mathbf{0} & \mathbf{0} & \mathbf{0} & \dots & \mathbf{x}^{(i)} \end{bmatrix} \begin{bmatrix} a_0^{(i)}(1-a_0^{(i)}) & -a_1^{(i)}a_0^{(i)} & -a_2^{(i)}a_0^{(i)} & \dots & -a_9^{(i)}a_0^{(i)} \\ a_0^{(i)}(1-a_1) & a_1^{(i)}(1-a_1) & -a_2^{(i)}a_1^{(i)} & \dots & -a_9^{(i)}a_1^{(i)} \\ a_0^{(i)}a_2^{(i)} & -a_1^{(i)}a_2^{(i)} & a_2^{(i)}(1-a_2^{(i)}) & \dots & -a_9^{(i)}a_2^{(i)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ -a_0^{(i)}a_9^{(i)} & -a_1^{(i)}a_9^{(i)} & -a_2^{(i)}a_9^{(i)} & \dots & -a_9^{(i)}(1-a_9^{(i)}) \end{bmatrix} \begin{bmatrix} -y_0^{(i)}/\hat{y}_0^{(i)} \\ -y_1^{(i)}/\hat{y}_1^{(i)} \\ -y_2^{(i)}/\hat{y}_2^{(i)} \\ \vdots \\ -y_9^{(i)}/\hat{y}_9^{(i)} \end{bmatrix} \right\}$$

$$+ \sum_{i=0}^{b-1} \begin{bmatrix} a_0^{(i)}(1-a_0^{(i)}) & -a_1^{(i)}a_0^{(i)} & -a_2^{(i)}a_0^{(i)} & \dots & -a_9^{(i)}a_0^{(i)} \\ a_0^{(i)}a_1^{(i)} & a_1^{(i)}(1-a_1) & -a_2^{(i)}a_1^{(i)} & \dots & -a_9^{(i)}a_1^{(i)} \\ a_0^{(i)}a_2^{(i)} & -a_1^{(i)}a_2^{(i)} & a_2^{(i)}(1-a_2^{(i)}) & \dots & -a_9^{(i)}a_2^{(i)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ -a_0^{(i)}a_9^{(i)} & -a_1^{(i)}a_9^{(i)} & -a_2^{(i)}a_9^{(i)} & \dots & -a_9^{(i)}(1-a_9^{(i)}) \end{bmatrix} \times \begin{bmatrix} -a_0^{(i)}/\hat{y}_0^{(i)} \\ -a_1^{(i)}/\hat{y}_1^{(i)} \\ -a_2^{(i)}/\hat{y}_2^{(i)} \\ \vdots \\ -a_9^{(i)}/\hat{y}_9^{(i)} \end{bmatrix} \mathbf{x}^{(i)T}$$

CCE loss and its gradient for the batch samples:

$$L = \frac{1}{b} [L_0 + \dots + L_{b-1}]$$

$$= \frac{1}{b} [-\mathbf{y}^{(0)T} \log(\hat{\mathbf{y}}^{(0)}) + \dots + -\mathbf{y}^{(b-1)T} \log(\hat{\mathbf{y}}^{(b-1)})].$$

$$[\nabla_{\hat{\mathbf{y}}^{(0)}}(L_0) \quad \dots \quad \nabla_{\hat{\mathbf{y}}^{(b-1)}}(L_{b-1})] = \begin{bmatrix} -y_0^{(0)}/\hat{y}_0^{(0)} & \dots & -y_0^{(0)}/\hat{y}_9^{(b-1)} \\ -y_1^{(0)}/\hat{y}_1^{(0)} & \dots & -y_1^{(b-1)}/\hat{y}_1^{(b-1)} \\ -y_2^{(0)}/\hat{y}_2^{(0)} & \dots & -y_2^{(b-1)}/\hat{y}_2^{(b-1)} \\ \vdots & & \vdots \\ -y_9^{(0)}/\hat{y}_9^{(0)} & \dots & -y_9^{(b-1)}/\hat{y}_9^{(b-1)} \end{bmatrix}$$

```
In [ ]: ## Define the loss function and its gradient
def cce(Y, Yhat):
    return np.mean(np.sum(-Y*np.log(Yhat), axis = 0))

def cce_gradient(Y, Yhat):
    return (-Y/Yhat)

# TensorFlow in-built function for categorical crossentropy loss
#cce = tf.keras.Losses.CategoricalCrossentropy()
```

Softmax activation layer class:

**Forward:**

$$[\mathbf{a}^{(0)} \quad \dots \quad \mathbf{a}^{(b-1)}] = [\text{softmax}(\mathbf{z}^{(0)}) \quad \dots \quad \text{softmax}(\mathbf{z}^{(b-1)})]$$

$$\Rightarrow \mathbf{A} = \text{softmax}(\mathbf{Z}).$$

**Backward:**

$$[\nabla_{\mathbf{z}^{(0)}}(L_0) \quad \dots \quad \nabla_{\mathbf{z}^{(b-1)}}(L_{b-1})] = [\nabla_{\mathbf{z}^{(0)}}(\mathbf{a}^{(0)}) \times \nabla_{\mathbf{a}^{(0)}}(L_0) \quad \dots \quad \nabla_{\mathbf{z}^{(b-1)}}(\mathbf{a}^{(b-1)}) \times \nabla_{\mathbf{a}^{(b-1)}}(L_{b-1})]$$

$$= \left[ \begin{bmatrix} a_0^{(0)}(1-a_0^{(0)}) & -a_1^{(0)}a_0^{(0)} & -a_2^{(0)}a_0^{(0)} & \dots & -a_9^{(0)}a_0^{(0)} \\ a_0^{(0)}a_1^{(0)} & a_1^{(0)}(1-a_1) & -a_2^{(0)}a_1^{(0)} & \dots & -a_9^{(0)}a_1^{(0)} \\ a_0^{(0)}a_2^{(0)} & -a_1^{(0)}a_2^{(0)} & a_2^{(0)}(1-a_2^{(0)}) & \dots & -a_9^{(0)}a_2^{(0)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ -a_0^{(0)}a_9^{(0)} & -a_1^{(0)}a_9^{(0)} & -a_2^{(0)}a_9^{(0)} & \dots & -a_9^{(0)}(1-a_9^{(0)}) \end{bmatrix} \times \begin{bmatrix} -y_0^{(0)}/\hat{y}_0^{(0)} \\ -y_1^{(0)}/\hat{y}_1^{(0)} \\ -y_2^{(0)}/\hat{y}_2^{(0)} \\ \vdots \\ -y_9^{(0)}/\hat{y}_9^{(0)} \end{bmatrix} \right] \dots \dots \dots \left[ \begin{bmatrix} a_0^{(b-1)}(1-a_0^{(b-1)}) & -a_1^{(b-1)}a_0^{(b-1)} & -a_2^{(b-1)}a_0^{(b-1)} & \dots & -a_9^{(b-1)}a_0^{(b-1)} \\ a_0^{(b-1)}a_1^{(b-1)} & a_1^{(b-1)}(1-a_1) & -a_2^{(b-1)}a_1^{(b-1)} & \dots & -a_9^{(b-1)}a_1^{(b-1)} \\ a_0^{(b-1)}a_2^{(b-1)} & -a_1^{(b-1)}a_2^{(b-1)} & a_2^{(b-1)}(1-a_2^{(b-1)}) & \dots & -a_9^{(b-1)}a_2^{(b-1)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ -a_0^{(b-1)}a_9^{(b-1)} & -a_1^{(b-1)}a_9^{(b-1)} & -a_2^{(b-1)}a_9^{(b-1)} & \dots & -a_9^{(b-1)}(1-a_9^{(b-1)}) \end{bmatrix} \times \begin{bmatrix} -y_0^{(b-1)}/\hat{y}_0^{(b-1)} \\ -y_1^{(b-1)}/\hat{y}_1^{(b-1)} \\ -y_2^{(b-1)}/\hat{y}_2^{(b-1)} \\ \vdots \\ -y_9^{(b-1)}/\hat{y}_9^{(b-1)} \end{bmatrix} \right]$$

```
In [ ]: '''Z = np.array([[1., 2., 3.], [4., 5., 6.]])
print(Z)
tf.nn.softmax(Z, axis = 0).numpy()'''

Out[ ]: 'Z = np.array([[1., 2., 3.], [4., 5., 6.]])\nprint(Z)\ntf.nn.softmax(Z, axis = 0).numpy()'
```

```
In [ ]: ## Softmax activation layer class
class Softmax(Layer):
    def forward(self, input):
        self.input = input
        self.output = tf.nn.softmax(self.input, axis = 0).numpy()

    def backward(self, output_gradient, learning_rate = None):
        ## Following is the inefficient way of calculating the backward gradient
        softmax_gradient = np.empty((self.input.shape[0], output_gradient.shape[1]), dtype = np.float64)
        for b in range(softmax_gradient.shape[1]):
            softmax_gradient[:, b] = np.dot((np.identity(self.output.shape[0]) - self.output[:, b].reshape(-1,1).T) * self.output[:, b].reshape(-1,1), output_gradient[:, b])
        return(softmax_gradient)
        ## Following is the efficient of calculating the backward gradient
        #T = (np.transpose(np.identity(self.output.shape[0]) - np.atleast_2d(self.output).T[:, np.newaxis, :], (1, 2, 0)) * np.atleast_2d(self.output))
        #return(np.einsum('ijk, ik -> jk', T, output_gradient))
```

Dense layer class:

**Forward:**

$$[\mathbf{z}^{(0)} \quad \dots \quad \mathbf{z}^{(b-1)} \dots] = \mathbf{W} [\mathbf{z}^{(0)} \quad \dots \quad \mathbf{z}^{(b-1)} \dots]$$

$$= [\mathbf{W}\mathbf{z}^{(0)} \quad \dots \quad \mathbf{W}\mathbf{z}^{(b-1)}]$$

$$\Rightarrow \mathbf{Z} = \mathbf{W}\mathbf{X}.$$

**Backward:**

$$\nabla_{\mathbf{W}}(L) = \frac{1}{b} \left[ \nabla_{\mathbf{W}}(\mathbf{z}^{(0)}) \times \nabla_{\mathbf{z}^{(0)}}(L) + \dots + \nabla_{\mathbf{W}}(\mathbf{z}^{(b-1)}) \times \nabla_{\mathbf{z}^{(b-1)}}(L) \right]$$

$$= \frac{1}{b} \left[ \nabla_{\mathbf{z}^{(0)}}(L) \mathbf{x}^{(0)T} + \dots + \nabla_{\mathbf{z}^{(b-1)}}(L) \mathbf{x}^{(b-1)T} \right].$$

```
In [ ]: ## Dense Layer class
class Dense(Layer):
    def __init__(self, input_size, output_size):
        self.weights = np.empty((output_size, input_size+1)) # bias trick
        self.weights[:, :-1] = 0.01*np.random.randn(output_size, input_size)
        self.weights[:, -1] = 0.01 # set all bias values to the same nonzero constant

    def forward(self, input):
        self.input = np.vstack([input, np.ones((1, input.shape[1]))]) # bias trick
        self.output = np.dot(self.weights, self.input)

    def backward(self, output_gradient, learning_rate):
```

```

## Following is the inefficient way of calculating the backward gradient
dense_gradient = np.zeros((self.output.shape[0], self.input.shape[0]), dtype = np.float64)

for b in range(output_gradient.shape[1]):
    dense_gradient += np.dot(output_gradient[:, b].reshape(-1, 1), self.input[:, b].reshape(-1, 1).T)

dense_gradient = (1/output_gradient.shape[1])*dense_gradient
#dense_gradient = np.mean(dense_gradient, axis = 1)
## Following is the efficient way of calculating the backward gradient
#dense_gradient = (1/output_gradient.shape[1])*np.dot(np.atleast_2d(output_gradient), np.atleast_2d(self.input).T)
self.weights = self.weights + learning_rate * (-dense_gradient)

```

Function to generate sample indices for batch processing according to batch size

```

In [ ]: ## Function to generate sample indices for batch processing according to batch size
def generate_batch_indices(num_samples, batch_size):
    # Reorder sample indices
    reordered_sample_indices = np.random.choice(num_samples, num_samples, replace = False)
    # Generate batch indices for batch processing
    batch_indices = np.split(reordered_sample_indices, np.arange(batch_size, len(reordered_sample_indices), batch_size))
    return(batch_indices)

```

Example generation of batch indices

```

In [ ]: ## Example generation of batch indices
num_samples = 64
batch_size = 8
batch_indices = generate_batch_indices(num_samples, batch_size)
print(batch_indices)

```

Train the 0-layer neural network using batch training with batch size = 16

```

In [ ]: ## Train the 0-layer neural network using batch training with batch size = 16
learning_rate = 0.01 # Learning rate
batch_size = 200 # batch size
nepochs = 20 # number of epochs
loss_epoch = np.empty(nepochs, dtype = np.float32) # create empty array to store Losses over each epoch

# Neural network architecture
dlayer = Dense(num_features, num_labels) # define dense Layer
softmax = Softmax() # define softmax activation Layer

# Steps: run over each sample in the batch, calculate loss, gradient of loss,
# and update weights.

epoch = 0
while epoch < nepochs:
    batch_indices = generate_batch_indices(num_samples, batch_size)
    loss = 0
    for b in range(len(batch_indices)):
        dlayer.forward(X_train[:,batch_indices[b]]) # forward prop
        softmax.forward(dlayer.output) # Softmax activate
        loss += cce(Y_train[:,batch_indices[b]], softmax.output) # calculate Loss
    # Backward prop starts here
    grad = cce_gradient(Y_train[:,batch_indices[b]], softmax.output)
    grad = softmax.backward(grad)
    grad = dlayer.backward(grad, learning_rate)
    loss_epoch[epoch] = loss/len(batch_indices)
    print("Epoch %d: loss = %f"%(epoch+1, loss_epoch[epoch]))
    epoch = epoch + 1

```

```


Epoch 1: loss = 1.442402
Epoch 2: loss = 0.849283
Epoch 3: loss = 0.683426
Epoch 4: loss = 0.602557
Epoch 5: loss = 0.553313
Epoch 6: loss = 0.519562
Epoch 7: loss = 0.494687
Epoch 8: loss = 0.475379
Epoch 9: loss = 0.459935
Epoch 10: loss = 0.447142
Epoch 11: loss = 0.436330
Epoch 12: loss = 0.427059
Epoch 13: loss = 0.419007
Epoch 14: loss = 0.411885
Epoch 15: loss = 0.405593
Epoch 16: loss = 0.399913
Epoch 17: loss = 0.394804
Epoch 18: loss = 0.390118
Epoch 19: loss = 0.385812
Epoch 20: loss = 0.381864

```

```

In [ ]: ## Plot training loss as a function of epoch:
plt.plot(loss_epoch)
plt.xlabel('Epoch')
plt.ylabel('Loss value')
plt.show()

```



```

In [ ]: ## Accuracy on test set
dlayer.forward(X_test)
softmax.forward(dlayer.output)
ypred = np.argmax(softmax.output.T, axis = 1)
print(ypred)
ytrue = np.argmax(Y_test.T, axis = 1)
print(ytrue)
np.mean(ytrue == ypred)

```

```
[7 2 1 ... 4 5 6]  
[7 2 1 ... 4 5 6]
```

```
Out[ ]: 0.9035
```

```
In [ ]: def plot_misclassified_images(X_test, ytrue, ypred, num_images=5):  
    misclassified_indices = np.where(ytrue != ypred)[0]  
    random_misclassified_indices = np.random.choice(misclassified_indices, num_images, replace=False)  
  
    fig, axes = plt.subplots(1, num_images, figsize=(20, 5))  
    for i, index in enumerate(random_misclassified_indices):  
        axes[i].imshow(X_test[:, index].reshape(28, 28), cmap='gray')  
        axes[i].set_title(f"True: {ytrue[index]}, Predicted: {ypred[index]}")  
        axes[i].axis('off')  
    plt.tight_layout()  
    plt.show()  
  
plot_misclassified_images(X_test, ytrue, ypred)
```

